

# Real-time NBA Player, Game, and Season Analytics

Ben Aoki-Sherwood\*      Arjun Rao†

## 1 Project Goals

We will build a scalable, real-time web application that predicts NBA team rankings and playoff probabilities using statistical and machine learning models. The system will:

1. **Real-time data ingestion and updates:** Periodically collect current NBA game results and team and player statistics from the [ESPN API](#) and an [NBA.com client API](#), and re-run analytics based on these live data.
2. **Player-, team- and season-level stats prediction:** Use deep learning (LSTM) and mathematical ([SpringRank](#)) models trained on recent historical data to predict:
  - Game outcomes
  - Player outcomes:
    - Box scores
    - Playing time
    - Injury
  - Season stats:
    - End-of-season team rankings for each conference
    - Playoff qualification probabilities for each team
    - Expected win-loss records based on current rankings and team schedules
3. **In-game event prediction:** use Poisson processes to predict:
  - Real-time win probabilities during live games
  - Player scoring rates and game milestone probabilities (e.g., “Will Player X score 25+ points tonight?”)
  - Event likelihoods (3PM, rebounds, blocks, steals, turnovers, fouls)
  - Expected scoring in specific time intervals
  - Scoring runs (what is the probability that Team A goes on a 10-0 run?)
  - Non-homogeneous Poisson processes to model momentum and clutch performance
4. **Dashboard:** Dashboard interface allows users to view current rankings, game event, score, and season projections, and historical trends with minimal latency

---

\*<https://github.com/aoki-sherwoodb>

†<https://github.com/arjunrao619>

## 2 Software and Hardware Components

### 2.1 Cloud Platform: GCP

### 2.2 Data Layer

1. **Cloud SQL (PostgreSQL)** - Relational database for:
  - Historical game results and statistics (2010-present)
  - Team and player information
  - Training dataset storage
  - User query logs
2. **Google Cloud Storage (GCS)** - Object storage for:
  - Trained ML model weights and checkpoints
  - Large datasets and CSV exports
  - Backup and archival data
3. **Memorystore (Redis)** - In-memory cache for:
  - Current team rankings and standings
  - Recent API responses
  - Frequently accessed predictions
  - Session management

### 2.3 Compute Layer

4. **Google Kubernetes Engine (GKE)** - Container orchestration hosting:
  - Data Ingestion Service (Python FastAPI)
  - ML Training/Inference Service (Python, scikit-learn/TensorFlow)
  - Poisson Event Modeling Service (Python, scipy)
  - API Gateway Service (Node.js/Python)
  - Web Frontend Service (React.js served via nginx)
5. **Docker** - Containerization of all microservices

### 2.4 Message Queue Layer

6. **Cloud Pub/Sub** - Message queue for:
  - Distributing incoming game results to processing services
  - Triggering model retraining jobs
  - Coordinating between microservices

## 2.5 APIs

### 7. REST APIs - For:

- External data ingestion (ESPN, NBA Stats)
- Internal microservice communication
- Frontend-backend communication
- Public API endpoints for client applications

## 2.6 ML

### 8. Python 3.10+ with libraries:

- scikit-learn (baseline models: Random Forest, Logistic Regression)
- TensorFlow/Keras (optional neural network models)
- scipy.stats (Poisson distributions and statistical modeling)
- numpy (numerical computations and Monte Carlo simulations)
- pandas (data processing)
- requests (API calls)

## 2.7 Frontend

### 9. React.js - Single-page application for user dashboard

### 10. Chart.js/D3.js - Data visualization libraries

### 11. Material-UI - Component library

## 2.8 Monitoring and Logging

### 12. Cloud Logging - Centralized logging

### 13. Cloud Monitoring - Performance metrics and alerting

## 2.9 Hardware Resources (GCP)

- **GKE Cluster:** 3-5 n1-standard-2 nodes (2 vCPUs, 7.5GB RAM each)
- **Cloud SQL Instance:** db-n1-standard-1 (1 vCPU, 3.75GB RAM)
- **Memorystore Redis:** 1GB Basic Tier instance

## 3 Architecture

## 4 Microservices and Interactions

### 4.1 Data Flow and Component Interactions

#### 4.1.1 Data Ingestion

1. **Data Ingestion service** runs on a scheduled basis (e.g., every 30 minutes during the NBA season)
2. It makes REST API calls to external sources (ESPN API, NBA Stats API) to fetch:

- Recent game results
  - Updated team standings
  - Player performance statistics
3. The service validates and transforms the raw data into a standardized format
  4. Validated data is published to **Cloud Pub/Sub** topic `new-game-results`
  5. The ingestion service also writes raw data to **Cloud SQL** for historical record-keeping

#### 4.1.2 Model Training and Inference

1. **ML Training/Inference service** subscribes to the `new-game-results` topic
2. When new game data arrives:
  - The service updates its training dataset in **Cloud SQL**
  - If enough new data has accumulated (e.g., 1-4 new games or daily on days when games are played), trigger a retraining job
3. Model training process:
  - Fetches historical data from **Cloud SQL** (last 5 seasons)
  - Trains ensemble models (Random Forest, Logistic Regression)
  - Validates model performance using cross-validation
  - Saves model weights and metadata to **Google Cloud Storage**
4. After training, the service generates fresh predictions for all teams
5. Predictions are written to both:
  - **Cloud SQL** (persistent storage with timestamps)
  - **Memorystore Redis** (fast cache with 1-hour TTL)

#### 4.1.3 Poisson Process Modeling

1. **Poisson Event Modeling service** subscribes to both `new-game-results` and `live-game-events` topics
2. **Rate parameter estimation:**
  - Service queries **Cloud SQL** for play-by-play data
  - Calculates scoring rates for teams/players:
    - Points per minute (overall and situational)
    - Three-point attempt rates
    - Turnover rates
    - Foul rates
  - Adjusts rates based on:
    - Opponent defensive ratings
    - Home/away status

- Rest days
  - Recent form (last 5 games)
  - Stores computed rate parameters in **Redis** with 24-hour TTL
- 3. Live game predictions:**
- When a live game event occurs (score update, timeout, etc.):
    - Service fetches current game state (score, time remaining, possessions)
    - Retrieves team  $\lambda$  parameters from **Redis**
    - Runs Monte Carlo simulations (10,000 iterations):
      - \* Simulates remaining game using dual Poisson processes
      - \* One process for each team
      - \* Accounts for time remaining and possession count
    - Calculates win probability, expected final score, and confidence intervals
    - Publishes results to **prediction-requests** topic
    - Caches live probabilities in **Redis** (1-minute TTL for rapidly changing values)

- 4. Player event predictions:**
- Calculates probabilities for player milestones:
    - “Will Player X score 25+ points tonight?”
    - “Probability of 5+ three-pointers made”
    - “Risk of fouling out”
  - Uses player-specific rate parameters adjusted for matchup difficulty
  - Results stored in **Redis** and **Cloud SQL**

- 5. Model persistence:**
- Fitted Poisson rate parameters saved to **GCS** daily
  - Historical accuracy metrics logged to **Cloud SQL**
  - Enables backtesting and model improvement

#### 4.1.4 API Requests

1. User visits the web application frontend
2. **Frontend (React)** sends HTTPS requests to the **API Gateway Service**
3. API Gateway checks authentication/authorization
4. For season-end prediction requests:
  - Gateway first checks **Redis cache** for recent predictions
  - If cache hit: Returns cached data immediately
  - If cache miss: Queries **Cloud SQL** for latest predictions from ML Training Service
  - Query results are cached in Redis for future requests

5. For live game probability requests:
  - Gateway checks **Redis** for most recent Poisson-based win probabilities
  - If not cached or stale, sends request to **Poisson Event Modeling Service** via Pub/Sub
  - Service computes fresh probability using Monte Carlo simulation and caches result
  - Returns win probability, expected score, and confidence intervals
6. For player event predictions:
  - Gateway queries cached Poisson rate parameters from **Redis**
  - If cache miss, **Poisson Service** computes rates from play-by-play data in **Cloud SQL**
  - Returns probabilities for requested milestones (e.g., “25+ points tonight”)
7. For historical data requests:
  - Gateway queries **Cloud SQL** directly
  - Results may be cached based on query patterns
8. Responses are sent back to the frontend as JSON

#### 4.1.5 Real-time Updates

1. When critical events occur (playoff scenarios change, major upsets):
  - Ingestion Service publishes to `model-retrain-trigger` topic
  - ML Service receives trigger and prioritizes immediate retraining
  - Once new predictions are ready, a notification is published
  - Frontend can poll for updates or use WebSocket connections for live updates
2. During live games:
  - Ingestion Service publishes to `live-game-events` topic every 30-60 seconds
  - **Poisson Event Modeling Service** immediately recalculates win probabilities
  - Updated probabilities pushed to **Redis**
  - Frontend displays real-time probability charts that update as game progresses

#### 4.1.6 Containerized Deployment

1. All services are packaged as **Docker containers**
2. **GKE** manages container lifecycle:
  - Auto-scaling based on CPU/memory usage
  - Load balancing across multiple pod replicas
  - Health checks and automatic restarts
  - Rolling updates for zero-downtime deployments
3. Services communicate within the cluster using Kubernetes service discovery

#### 4.1.7 Message Queue Coordination

**Cloud Pub/Sub** decouples services and enables:

- Asynchronous processing (ingestion doesn't wait for training)
- Fault tolerance (messages are persisted if subscribers are down)
- Scalability (multiple service instances can subscribe to same topic)
- Priority handling (urgent retraining can use separate topic)

## 5 Debugging and Testing

### 5.1 Unit Testing

- **Component-level tests** for each microservice
- **Framework:** pytest for Python services, Jest for React frontend
- **Coverage targets:** ≥80% code coverage
- **Testing approach:**
  - Static external API calls using historical games from ESPN API
  - Verify ML model input/output shapes
  - Test cache hit/miss scenarios

### 5.2 Integration Testing

- **Service-to-service communication tests**
- **Pub/Sub message flow verification:**
  - Publish test messages and verify subscribers process them correctly
  - Test message acknowledgment and retry logic
- **Database interaction tests:**
  - Test CRUD operations against test PostgreSQL database
  - Verify transaction handling and rollback scenarios
- **Cache integration tests:**
  - Verify Redis caching behavior
  - Test cache invalidation logic

### 5.3 End-to-End Testing

- **Full pipeline tests:**
  - Simulate game data ingestion to model training to prediction generation to API response pipeline
  - Use staging environment with production-like configuration
- **Load testing:**
  - Apache JMeter or Locust to simulate concurrent users

- Test API Gateway under load (100+ requests/second)
- Verify auto-scaling behavior in GKE

## 5.4 Model Validation and Testing

### ML Models:

- Training data validation:
  - Check for data quality issues (missing values, clip outlier statistics)
  - Verify temporal consistency (no data leakage)
- Model performance metrics:
  - Accuracy, precision, recall for playoff predictions
  - Mean Absolute Error (MAE) for win-loss projections
  - Backtesting on previous seasons (2022-2024)
- A/B testing framework:
  - Train and deploy several model architectures and compare performance on both historical and current season data

### Poisson process model development and validation:

- Rate parameter validation:
  - Cross-validation: Train on games 1-70, test on games 71-82 of a season
  - Verify rate parameter estimates stabilize after enough games
  - Compare estimated rates to actual observed rates using Chi-square goodness-of-fit tests
- Win probability accuracy:
  - Calibration plots: Are predicted 70% win probabilities actually winning 70% of the time?
  - Brier score: Measure accuracy of probabilistic predictions
  - Compare to baseline models (Vegas odds, SpringRank)
  - Track accuracy across different game situations (close games, blowouts, clutch time)
- Monte Carlo simulation validation:
  - Verify convergence: Results stable with 10,000 simulations
  - Compare Monte Carlo to analytical solutions for simple cases
  - Test computational performance (simulations complete in ~real-time)
- Player event prediction validation:
  - Backtest player milestone predictions on historical games
  - Calculate ROI if used for betting (theoretical analysis only)
  - Verify confidence intervals contain true outcomes at expected rates
- Assumption checks:

- Test for independence: Are consecutive baskets independent? (probably not!)
- Test for constant rate: Do rates change during game? (we expect difference for 1st quarter vs 4th)
- Implement piecewise Poisson models if needed (different  $\lambda$  for different game periods)

- **Edge cases:**

- Extreme scores (50-point leads, overtime games)
- Unusual rates (defensive battles, high-scoring shootouts)

## 5.5 Debugging Tools and Techniques

- **Cloud logging:**

- Structured logging with correlation IDs across services
- Log levels: DEBUG for development, INFO for production
- Searchable logs aggregated in Google Cloud Console

- **Cloud monitoring:**

- Custom metrics for:
  - \* API response times
  - \* ML model training duration
  - \* Poisson simulation computation time
  - \* Prediction accuracy over time (both ML and Poisson models)
  - \* Cache hit rates (especially for Poisson rate parameters)
  - \* Live game probability update frequency
- Alerts for:
  - \* Error rate spikes
  - \* API latency  $> 500\text{ms}$
  - \* Failed model training jobs
  - \* Poisson simulations taking  $> 200\text{ms}$
  - \* Win probability calculation errors
  - \* Unusual rate parameter values (potential data quality issues)

- **Distributed tracing:**

- Cloud Trace to track request flow across microservices
- Identify bottlenecks in the request pipeline

- **Local development:**

- Docker Compose for running full stack locally
- Minikube for local Kubernetes testing
- Environment parity (dev/staging/prod)

## 5.6 Data Quality Monitoring

- **Ingestion monitoring:**
  - Track API availability and response times
  - Alert on data schema changes
  - Validate data completeness (all expected games ingested)
- **Anomaly detection:**
  - Flag unusual prediction patterns
  - Detect model drift (degrading accuracy over time)

# 6 Cloud Technologies Used

## 6.1 Message Queues

- **Technology Used:** Google Cloud Pub/Sub
- **For:** Asynchronous message passing between Data Ingestion Service, ML Training Service, Poisson Process Service, and API Gateway

## 6.2 Databases

- **Technology Used:** Cloud SQL (PostgreSQL)
- **For:** Persistent storage of historical game data, team statistics, player information, and prediction history

## 6.3 Key-Value Stores

- **Technology Used:** Memorystore (Redis)
- **For:** High-speed caching of current rankings, recent predictions, and API responses

## 6.4 Containerization

- **Technology Used:** Docker containers orchestrated by Google Kubernetes Engine (GKE)
- **For:** Containerization and orchestration of all microservices (ML inference/training, Poisson process modeling, API gateway)

## 6.5 Storage Services

- **Technology Used:** Google Cloud Storage (GCS)
- **For:** Object storage for trained ML model weights, checkpoints, and large datasets

## 6.6 RPC / API Interfaces

- **Technology Used:** REST APIs (internal and external)
- **For:** Communication between frontend and backend, between microservices, and with external data providers

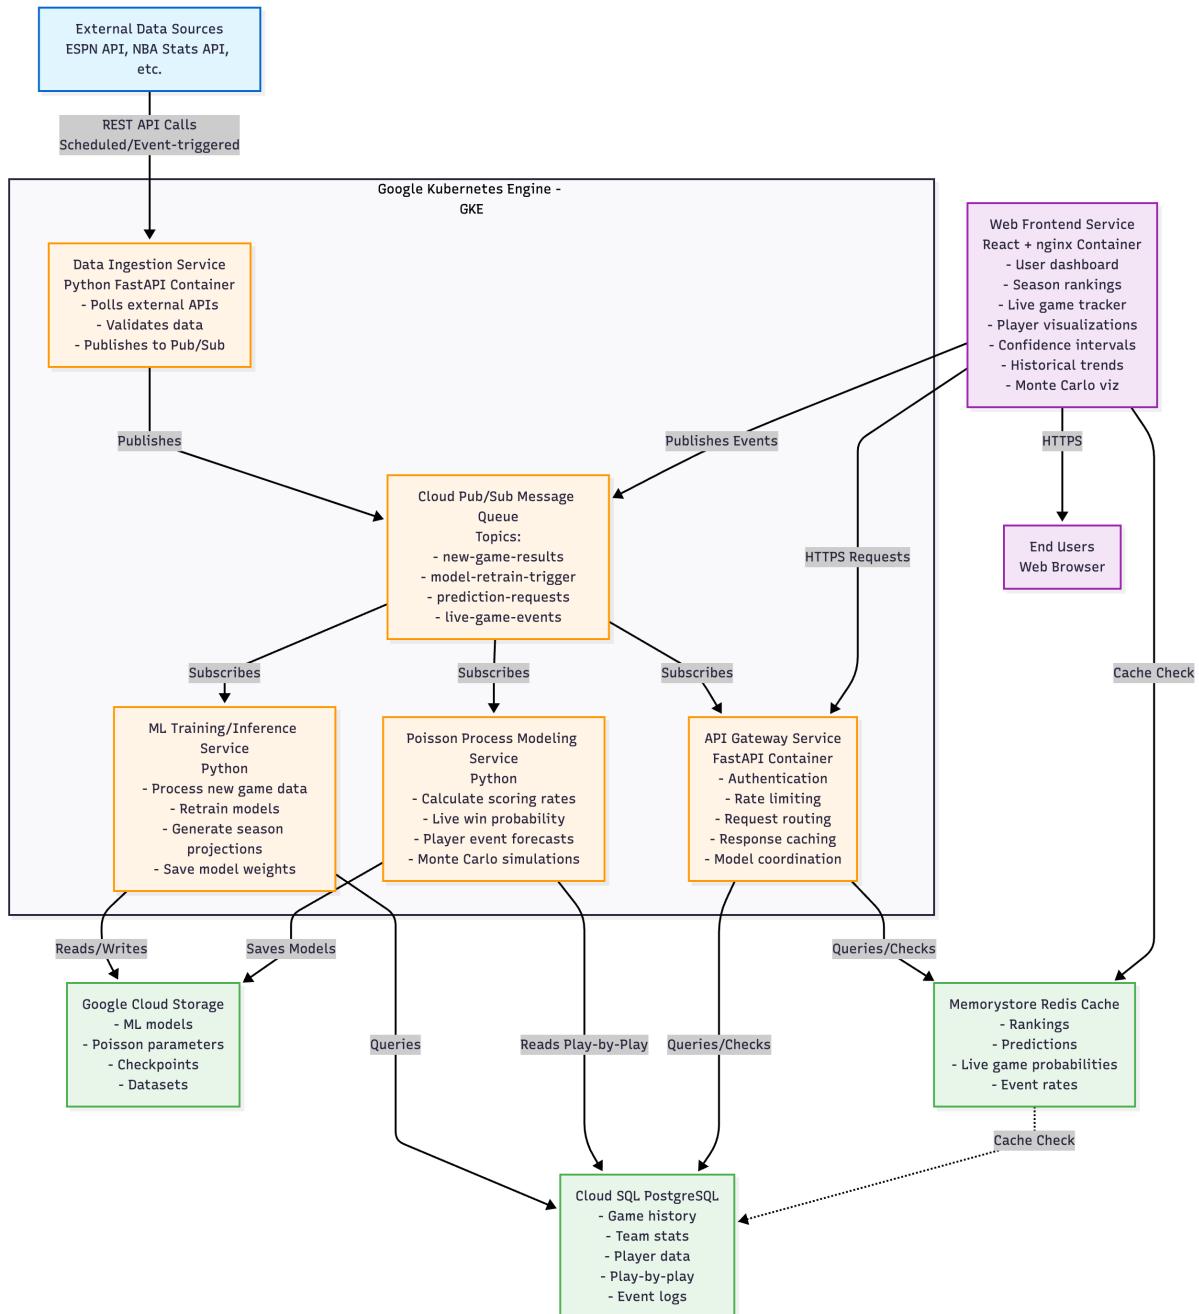


Figure 1: System Architecture Diagram